

## EE-411, HomeWork 3 : Neural networks

This homework involves some coding, with the language of your choice. Present your results with graphs, plots, and data. Jupyter notebooks are a good option, and we recommend you to send your work as a notebook on Google colab.

### 1 Backpropagation with logistic loss

Let  $f_{\mathbf{w}} : \mathbb{R}^D \rightarrow \mathbb{R}$  be a single-hidden-layer, fully connected, neural network with no biases and parameters  $\mathbf{w} = \{\mathbf{W}^{(1)}, \mathbf{w}^{(2)}\}$ , such that its forward pass computes

$$\mathbf{x}^{(1)} = \sigma(\mathbf{z}^{(1)}) = \sigma(\mathbf{W}^{(1)}\mathbf{x}^{(0)}) \quad (1)$$

$$\hat{y} = \sigma(z^{(2)}) = \sigma(\mathbf{w}^{(2)\top} \mathbf{x}^{(1)}). \quad (2)$$

Here,  $\sigma(\cdot)$  denotes a sigmoid activation. In what follows, we will assume  $D = 5$  and  $K = 6$ .

1. Implement a function `predict(X,W)` that given a batch of  $B$  samples  $\{\mathbf{x}_i\}_{i=1}^B$  (an array `X` of size  $B \times D$ ) and a dictionary of weights (`W={w_1: D x K, w_2: K x 1}`) computes the set of outputs  $\{(\mathbf{z}_i^{(1)}, z_i^{(2)}, \hat{y}_i)\}_{i=1}^B$ . *Hint* : Try to avoid using `for` loops.
2. Implement a function `logistic_loss(y, y_hat)` that given a vector of labels  $\mathbf{y} \in \{0,1\}^B$  and a vector of predictions  $\hat{\mathbf{y}} \in \mathbb{R}^B$  computes the average logistic loss of a batch. Compute the average loss for the case  $\hat{\mathbf{y}} = \mathbf{0}$  and  $\mathbf{y} = \mathbf{0}$ . Explain your result.
3. Implement another function `stable_logistic_loss(y, z_2)` that given a vector of labels  $\mathbf{y} \in \{0,1\}^B$  and a vector of logits  $\mathbf{z}^{(2)} \in \mathbb{R}^B$  computes the average logistic loss of a batch. Make sure that your function is stable. Compute the average loss for the case when  $\mathbf{z}_2 = -10^{10} \cdot \mathbf{1}$  and  $\mathbf{y} = \mathbf{0}$ . Explain your result. *Hint* : You can use `np.logaddexp`.
4. Derive analytically, using the backpropagation algorithm, the expressions for the partial derivatives of the loss with respect to the weights, i.e.,

$$\frac{\partial \mathcal{L}(\mathbf{x}, y; \mathbf{w})}{\partial w_{ij}^{(1)}} \quad \text{and} \quad \frac{\partial \mathcal{L}(\mathbf{x}, y; \mathbf{w})}{\partial w_i^{(2)}}, \quad (3)$$

when the loss is computed using a single pair  $(\mathbf{x}, y)$ . Recall that we are using the logistic loss, and that the final implementation should be stable.

5. Implement a function `gradient(X,y,W)` which computes the gradient (i.e., vector of partial derivatives) of the average loss with respect to all the weights for a batch  $\{(\mathbf{x}_i, y_i)\}_{i=1}^B$ . *Hint* : Try to vectorize your code and expressions.

### 2 Classifying FashionMNIST using neural networks

In this exercise you will play with the dataset called **FashionMNIST**. Some examples are shown in Figure 1. You can use code snippets from the labs.

1. Load the dataset, construct the dataloaders and visualize the data. Which transform is necessary<sup>1</sup>?
2. **MultiLayer Perceptron (MLP)** : Construct a two hidden-layer MLP with 100 neurons and ReLU activation functions to classify **FashionMNIST**. Train this model using the cross-entropy loss using SGD (`lr=0.01`), SGD with momentum (`lr=0.01`, `momentum=0.9`, `nesterov=True`) and Adam (`lr=0.1`), and plot the training and test learning curves for every epoch on a single plot. Comment on your results.

---

1. for the samples to be compatible with the models we will create



FIGURE 1 – Some examples from the FashionMNIST dataset.

3. **Convolutional Neural Network (CNN)** : Construct a CNN with three convolutional layers (`kernel_size=3`) and 16, 32, and 64 channels, respectively; one max-pooling layer (`kernel_size=2`) after every convolution; and a final fully connected layer. Train this model with the same configurations as before, and plot the training and validation learning curves on a single plot. Comment on your results.
4. Create a function that computes the number of parameters of a given model. Show the number of parameters for the two models you have used. Does more parameters translate to better performance? Explain.
5. **PermutedFashionMNIST<sup>2</sup>** : In this version of the dataset, the pixels are randomly permuted. Visualize the new dataset. Train one MLP and one CNN with the best parameters from before. What do you observe? Is one model more affected than the other? Explain.
6. (*Bonus*) Using any of the tips and tricks seen in class or during the exercise sessions, optimize the test performance on the original FashionMNIST. Also feel free to explore things beyond the ones seen in class. In any case, explain your decisions. Did you manage to improve the performance? Why?

```
class RandomPermutation(object):
    def __init__(self, num_features):
        self.num_features = num_features
        self.reindex = torch.randperm(num_features)

    def __call__(self, img):
        assert self.num_features == img.numel()

        orig_shape = img.shape
        img = img.view(-1)[self.reindex].view(orig_shape)

        return img
```

---

2. Add the tranformation `RandomPermutation` to the list of transforms. This transform should be the last one.