EE-411, HomeWork 1: Maximum Likelihood & Probability

This homework involves some coding, with the language of your choice. Present your results with graphs, plots, and data. Jupyter notebooks are a good option, and we recommend you to send your work as a notebook on Google colab.

1 Statistical inference and Maximum Likelihood

$$P_{\lambda}(x_i) = \begin{cases} \frac{e^{-x_i/\lambda}}{Z(\lambda)} & \text{if } 1 \le x_i \le 20\\ 0 & \text{otherwise} \end{cases}$$
 (1)

We shall assume we only know that, for the true λ (denoted λ^*) $50 > \lambda^* > 0.05$.

- 1. Compute $Z(\lambda)$ such that the probability in eq.(1) is normalized. $\{x_i\}_{i=1}^n$? Show that $\mathbb{E}[X] = \lambda^2 \partial_{\lambda} \log Z(\lambda)$ and $\operatorname{var}[X] = \lambda^2 \partial_{\lambda} \mathbb{E}[X]$. Compute and plot the mean and variance as a function of λ .
- 2. What is the probability $P_{\lambda}(\{x_i\}_{i=1}^n)$ to observe a set of n events at position $\{x_i\}_{i=1}^n$? What is the log-likelihood function $\mathcal{L}(\lambda, \{x_i\}_{i=1}^n)) = \log(P_{\lambda}(\{x_i\}_{i=1}^n)/n$?
- 3. We can write a program that simulates n such observations sampled from the probability distribution (1) for a fixed λ . This can be done, for instance, with the numpy.random.exponential command in python, keeping only events x that are in [1,20]. We choose the true value to be $\lambda^* = 10$. Generate n = 10 observations. Plot the likelihood as a function of λ . Repeat for n = 20,100 and discuss what you see.
- 4. We now assume that we are given a set of n observations $\{x_i\}_{i=1}^n$, without being told the true value of λ^* . We consider the maximum likelihood estimator

$$\hat{\lambda}_{\mathrm{ML}}(\{x_i\}_{i=1}^n) = \operatorname{argmax}_{\lambda} \log \left(P_{\lambda}(\{x_i\}_{i=1}^n) \right) \tag{2}$$

and we shall define the squared error as $SE = (\hat{\lambda}_{ML}(\{x\}) - \lambda^*)^2$.

Create some data set with n = 10, 100, 1000 for different values of λ^* between 0 and 50 and see how the ML estimator performs. Note that finding the maximizer $\hat{\lambda}_{ML}$ can be done numerically in python, for instance using SciPy.

5. Show that the Fisher score information in this problem is given by

$$I(\lambda) = \mathbb{E}\left[\left(\frac{\partial}{\partial \lambda} \log P_{\lambda}(X)\right)^{2}\right] = \frac{1}{\lambda^{4}} \mathbb{E}\left[\left(X - \mathbb{E}[X]\right)^{2}\right] = \frac{\operatorname{var}[X]}{\lambda^{4}}$$
(3)

.

6. Another interesting estimator is given by maximum as posteriori (MAP) with the Jeffreys prior :

$$\hat{\lambda}_{J}(\{x_{i}\}_{i=1}^{n}) = \operatorname{argmax}_{\lambda} \log \left(P_{\lambda}(\{x_{i}\}_{i=1}^{n}) \sqrt{I(\lambda)} \right)$$
(4)

where $I(\lambda)$ is the Fisher information. Implement this estimator using SciPy.

- 7. If we average of many realizations (say about a hundred) we can obtain numerically the averaged mean squared error $MSE(\lambda^*, \hat{\lambda}, n)$ which is thus a function of n, λ^* and of the estimator $\hat{\lambda}$. Compute and plot, for many values of n from n = 10 to n = 1000, the curves $MSE(\lambda^*, \hat{\lambda}_{ML}, n)$ and $MSE(\lambda^*, \hat{\lambda}_{J}, n)$ as a function of λ^* .
- 8. How does the MSE curves at various n e compare with the Cramér-Rao bound for unbiased estimator $\text{MSE}(\hat{\lambda}) \geq \frac{1}{nI(\lambda^*)}$ (where $I(\lambda)$ is the Fisher information)? How does the Jeffrey and ML estimator behave? Which one would you choose?

2 Probability bounds and a pooling problem

We are going to follow the steps we took in lecture 1 and prove an interesting inequality: Let Z_1, \ldots, Z_m be independent random variables such that $Z_i = 1$ with probability p, and 0 with probability 1 - p. Then, for any $\epsilon \ge 0$ we have

$$\mathbb{P}\left(\frac{1}{m}\sum_{i}Z_{i} \ge p + \epsilon\right) \le e^{-2m\epsilon^{2}} \tag{5}$$

1. Our starting point is to realize that $\mathbb{P}(a \geq b) = \mathbb{P}(e^{\lambda a} \geq e^{\lambda b})$ for any $\lambda \geq 0$. Using Markov inequality and the proof strategy discussed in lecture 1, show that :

$$\mathbb{P}\left(\frac{1}{m}\sum_{i} Z_{i} \geq p + \epsilon\right) \leq \left(\frac{pe^{\lambda} + (1-p)}{e^{\lambda(p+\epsilon)}}\right)^{m}$$

2. Using the value of λ that minimizes the right-hand side of the former equation, show that

$$\mathbb{P}\left(\frac{1}{m}\sum_{i} Z_{i} \ge p + \epsilon\right) \le e^{-mf(p,\epsilon)}$$

with

$$f(p,\epsilon) = -\log\left(\left(\frac{p}{p+\epsilon}\right)^{p+\epsilon} \left(\frac{1-p}{1-p-\epsilon}\right)^{1-p-\epsilon}\right)$$

3. Show that

$$f(p, \epsilon = 0) = 0$$
, $\frac{\partial f(p, \epsilon)}{\partial \epsilon}\Big|_{\epsilon = 0} = 0$, and that $\frac{\partial^2 f(p, \epsilon)}{\partial \epsilon^2} \le 4$ for any ϵ .

4. Use Taylor's theorem (that states that $f(p,\epsilon) = f(p,0) + \epsilon f'(p,0) + \epsilon^2 f''(p,\tilde{\epsilon})/2$ for some unknown $\tilde{\epsilon}$, and where the prime stands for derivative with respect to ϵ) to show that $f(p,\epsilon) \leq 2\epsilon^2$, and prove the inequality (5).

Similarly, it is possible to show (Bonus 1: prove it) that

$$\mathbb{P}\left(\frac{1}{m}\sum_{i}Z_{i} \le p - \epsilon\right) \le e^{-2m\epsilon^{2}} \tag{6}$$

so that

$$\mathbb{P}\left(\left|\frac{1}{m}\sum_{i}Z_{i}-p\right|\geq\epsilon\right)\leq2e^{-2m\epsilon^{2}}\tag{7}$$

- 5. The number one use of such a bound is in terms of pooling problems. Suppose you want to know what fraction of the population in a country approves its current president: how many people should you ask to be confident, with probability at least 95 percent, that the error in estimating the fraction of people who approves the president is correct within one percent (so that \hat{p} is in [p-0.01, p+0.01] with 95% probability)?
- 6. Compare the number m^* you find this way with what you observe when performing numerical experiments in python:
 - Define a function that takes m and the true p as arguments and returns a random array of m votes (it has to work with generic m and p).
 - Starting with fixed values of $m = m^*$ and $p \in \{0.2, 0.5, 0.8\}$
 - Use this function evaluated in m^* and p to simulate polls.
 - Just by using the generated votes, estimate p.
 - Quantify the probability that \hat{p} is correct within one percent.
 - Which values of p seem to be harder to estimate? Do you find that the bound is accurate, or does it grossly overestimate the needed number?
 - For each p, repeat for different values of m to find the value that (more or less) gives an estimate which is correct within one percent with 95% probability.
 - Bonus 2: Plot the behaviour of the probability of error $\mathbb{P}(\hat{p} \notin [p-0.01, p+0.01])$ as a function of p for values of $m \in [10, 10^4]$, and compare it with the theoretical m^* .